

计算机科学与技术学院神经网络与深度学习课程实验报告

实验题目: image classification pipeline,		学号: 201900130024
日期: 2021. 10. 6	班级: 数据 19	姓名: 刘士渤
Email: liuburger@qq.com		
<p>实验目的:</p> <ul style="list-style-type: none">• understand the basic Image Classification pipeline and the data-driven approach (train/predict stages)• understand the train/val/test splits and the use of validation data for hyperparameter tuning.• develop proficiency in writing efficient vectorized code with numpy• implement and apply a k-Nearest Neighbor (kNN) classifier• implement and apply a Multiclass Support Vector Machine (SVM) classifier• implement and apply a Softmax classifier• implement and apply a Three layer neural network classifier• understand the differences and tradeoffs between these classifiers• get a basic understanding of performance improvements from using higher-level representations than raw pixels (e.g. color histograms, Histogram of Gradient (HOG) features)		
<p>实验软件和硬件环境:</p> <p>vscode jupyternotebook</p> <p>联想拯救者 Y7000p</p>		
<p>实验原理和方法:</p> <p>knn、svm、softmax、three layer net</p>		
<p>实验步骤: (不要求罗列完整源代码)</p> <p>一. knn</p> <p>1. 完成 k_nearest_neighbor.py</p> <p>两重循环非常简单, 对应位置做点积再开平方就行:</p> <pre>dists[i][j]=np.dot((X[i]-self.X_train[j]),(X[i]-self.X_train[j]))**0.5</pre> <p>一重循环利用 np. sum 省去了一重循环:</p> <pre>dis=np.sqrt(np.sum(np.square(X[i]-self.X_train),axis=1)) dists[i,:]=dis</pre> <p>无循环需要另辟蹊径, 将欧氏距离展开:</p> <pre>XY=np.dot(X,self.X_train.T) Xi=np.sum(np.square(X),axis=1).reshape([num_test,1])#按行加 Yj=np.sum(np.square(self.X_train),axis=1).reshape([num_train,1]) dists=np.sqrt(Xi-2*XY+np.array(Yj).T)</pre>		

根据参数 k ，使用大多数投票的方法确定预测标签


```
tmp=np.argsort(dists[i])
if k==1:
    y_pred[i]=self.y_train[tmp[0]]
elif k>1:
    for j in range(k):
        closest_y.append(self.y_train[tmp[j]])
    y_num={}
    for label in closest_y:
        if y_num.get(label)==None:
            y_num[label]=1
        else:
            y_num[label]+=1
    y_num=sorted(y_num.items(),key=lambda x:x[1],reverse=True)
    y_pred[i]=y_num[0][0]
```

2. 完成 knn. ipynb

在 k 的集合中用 5 折交叉验证选出最合适的 k :

```
for k in k_choices:
    k_to_accuracies[k]=[]
    for i in range(0,num_folds):
        classifier.train(tmpx[i],tmpy[i])
        dis=classifier.compute_distances_no_loops(X_train_folds[i])
        y_pred = np.array(classifier.predict_labels(dis, k=k))
        num_correct = np.sum(y_pred == y_train_folds[i])
        accuracy = float(num_correct) / len(y_pred)
        k_to_accuracies[k].append(accuracy)
```

发现在 $k=10$ 的时候效果最好:



用最好的 k 跑，结果达到了 28%

```
Got 144 / 500 correct => accuracy: 0.288000
```

二. svm

1. 完成 linear_classifier

采用 SGD，所以要有随即成分、梯度下降（-gradient）

```
randi=np.random.choice(num_train,batch_size)
X_batch=X[randi]
y_batch=y[randi]

self.W-= learning_rate*grad
```

预测就是选择使目标函数最大的集合

```
y_pred=np.argmax(np.dot(X,self.W),axis=1)
```

2. 完成 linear_svm

naïve 的关键部分人家已经写好了，自己只需要给 dW 加上正则项的导数：

```
dW+=2*reg*W
```


vectorized 复杂一些：

```
scores=X.dot(W)
#correct_scores的每一行都是正确的分数
correct_scores=scores[np.arange(num_train),y]
correct_scores=np.reshape(np.repeat(correct_scores,num_classes),(num_train,num_classes))
```

```
margin=scores+1.0-correct_scores
#对应相减的地方会是1, 则置零
margin[np.arange(num_train),y]=0
#margin现在是loss的初级，要得到loss只需要将正的项加起来
loss=(np.sum(margin[margin>0]))/num_train
loss+=reg*np.sum(W*W)
#分成两类
margin[margin>0]=1
margin[margin<=0]=0
row_sum = np.sum(margin, axis=1)# 1 by N
margin[np.arange(num_train), y] = -row_sum
dW += np.dot(X.T, margin)      # D by C#
#除以N再正则化
dW/=num_train
dW+=2*reg*W
```

3. 完成 svm.ipynb


绘制出平均图像:



比对 naïve 和 vectorized 的差异 (无差异):

```
Naive loss: 8.639771e+00 computed in 0.251944s
Vectorized loss: 8.639771e+00 computed in 0.004991s
difference: -0.000000
```

利用 SGD, 损失随迭代次数下降:



训练和预测准确率

```
training accuracy: 0.373612
validation accuracy: 0.378000
```


循环调参:

```
for rate in learning_rates:
    for strength in regularization_strengths:
        svm = LinearSVM()
        loss_hist = svm.train(X_train, y_train, learning_rate=rate, reg=strength,
                               num_iters=500, verbose=True)
        y_train_pred = svm.predict(X_train)
        y_val_pred = svm.predict(X_val)
        acc=np.mean(y_val == y_val_pred)
        if acc>best_val:
            best_val=acc
            best_svm=svm
        results[(rate,strength)]=(np.mean(y_train == y_train_pred),acc)
```

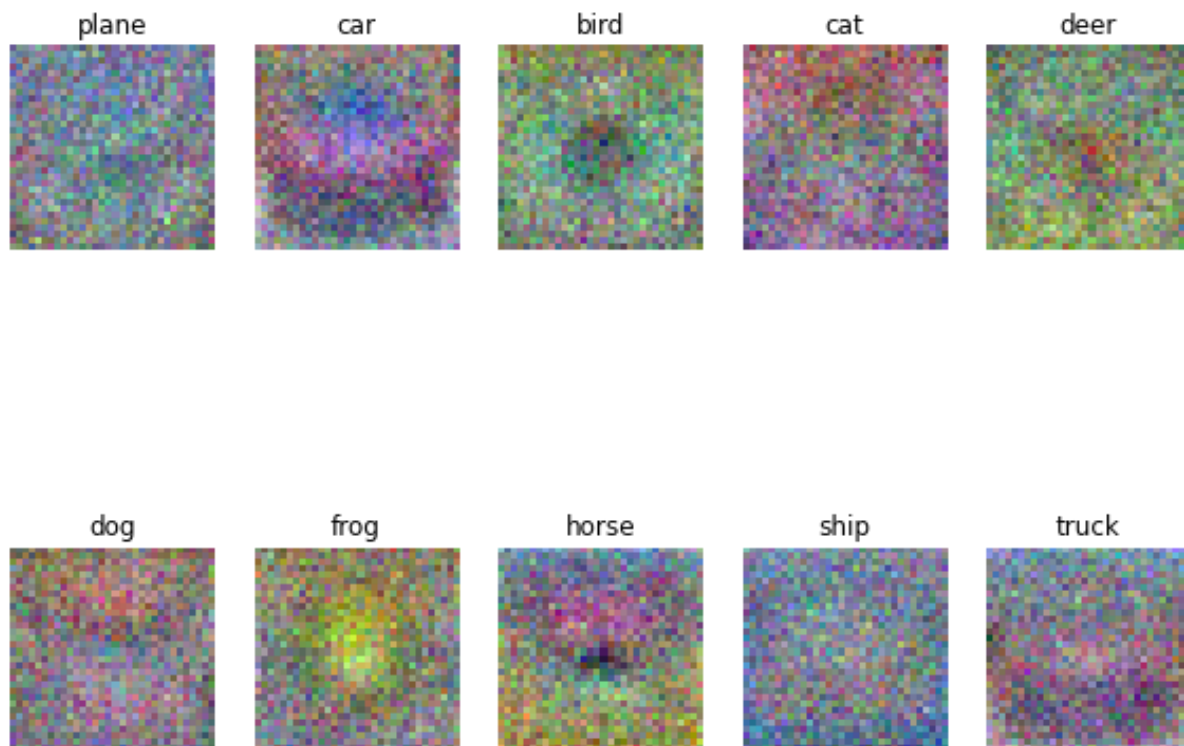
得到最佳的准确率:

best validation accuracy achieved during cross-validation: 0.398000

在数据集 CIFAR10 上的准确率:



模型可视化（汽车和青蛙还是可以隐约看出来）:



三. softmax

1. 完成 softmax.py

naïve 部分：求 f 函数、损失函数，求导。

```
N,C=X.shape[0],W.shape[1]
for i in range(N):
    f=np.dot(X[i],W)
    f-=np.max(f)#1*C
    s=np.sum(np.exp(f))
    loss+=np.log(s)-f[y[i]]
    dW[:,y[i]]-=X[i]
    for j in range(C):
        dW[:,j]+=np.exp(f[j])/s*X[i]

loss=loss/N+0.5*reg*np.sum(W*W)
dW=dW/N+reg*W
```

vectorized 部分：类似的方法，只是不用循环用矩阵的方法。

```
N=X.shape[0]
f=np.dot(X,W)#N*C
f-=f.max(axis=1).reshape(N,1)
s=np.exp(f).sum(axis=1)
loss=np.log(s).sum()-f[range(N),y].sum()

counts=np.exp(f)/s.reshape(N,1)
counts[range(N),y]-=1
dW=np.dot(X.T,counts)

loss=loss/N+0.5*reg*np.sum(W*W)
dW=dW/N+reg*W
```

2. 完成 softmax.ipynb

比较 naïve 和 vectorized 的 loss 发现无差别：

```
naive loss: 2.427571e+00 computed in 0.142212s
vectorized loss: 2.427571e+00 computed in 0.006065s
Loss difference: 0.000000
Gradient difference: 0.000000
```

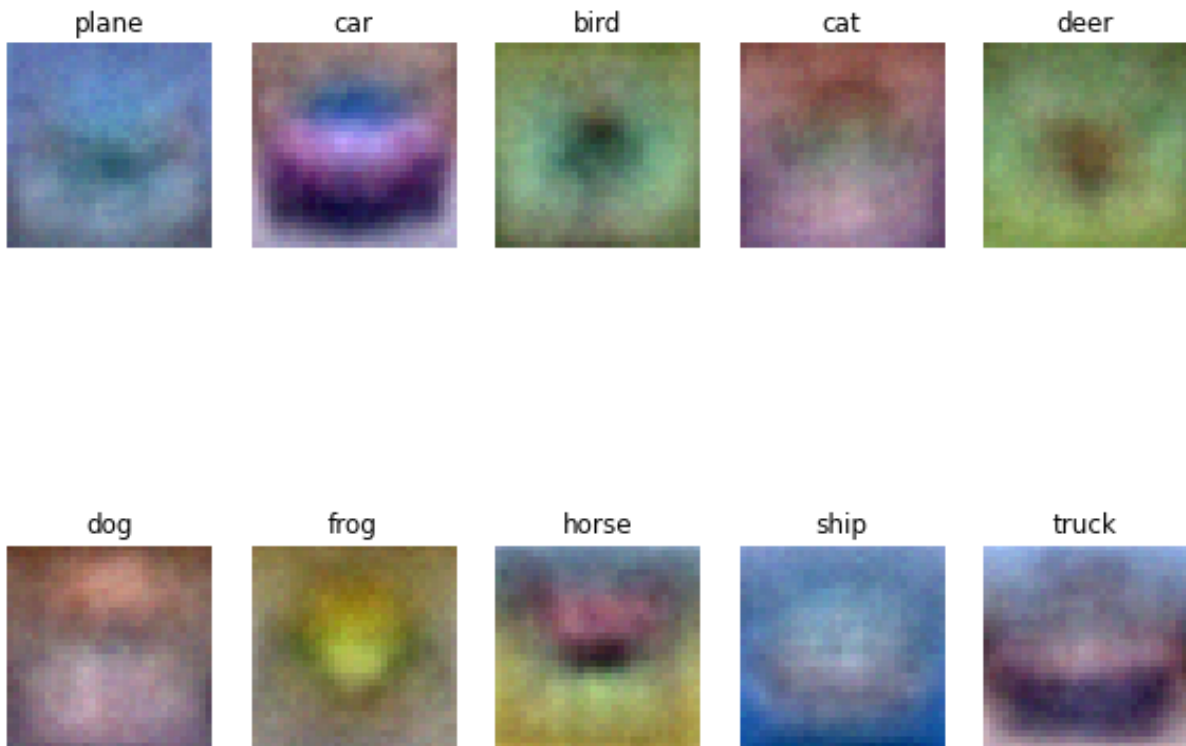
循环调参

```
for lr in learning_rates:
    for rs in regularization_strengths:
        softmax=Softmax()
        softmax.train(X_train,y_train,learning_rate=lr,reg=rs,num_iters=500,verbose=True)
        y_pred_train = softmax.predict(X_train)
        acc_train = np.mean(y_pred_train == y_train)
        y_pred_val = softmax.predict(X_val)
        acc_val = np.mean(y_pred_val == y_val)
        results[(lr, rs)] = (acc_train, acc_val)
        if acc_val > best_val:
            best_val = acc_val
            best_softmax = softmax
```

得到最佳的准确率:

best validation accuracy achieved during cross-validation: 0.353000

模型可视化 (效果要比 svm 好得多):



四. three layer net

1. 完成 neural_net.py

loss 函数:

hiddenlayer 就是 $y=WX+b$ 的形式

```
hidden_layer1=np.maximum(0,np.dot(X,W1)+b1)
hidden_layer2=np.maximum(0,np.dot(hidden_layer1,W2)+b2)
scores=np.dot(hidden_layer2,W3)+b3
```

f 函数和 loss 函数 (loss 要加正则项):

```
f=scores-np.max(scores,axis=1,keepdims=True)
loss=-f[range(N),y].sum()+np.log(np.exp(f).sum(axis=1)).sum()
loss=loss/N+0.5*reg*(np.sum(W1*W1)+np.sum(W2*W2)+np.sum(W3*W3))
```

根据 Computational graphs + Backpropagation 原理反向求导:

```
dscore=np.exp(f)/np.exp(f).sum(axis=1,keepdims=True)
dscore[range(N),y]-=1
dscore/=N
#反向传播
grads['W3']=np.dot(hidden_layer2.T,dscore)+reg*W3
grads['b3']=np.sum(dscore,axis=0)
dhidden2=np.dot(dscore,W3.T)
dhidden2[hidden_layer2<=0.00001]=0

grads['W2']=np.dot(hidden_layer1.T,dhidden2)+reg*W2
grads['b2']=np.sum(dhidden2,axis=0)
dhidden1=np.dot(dhidden2,W2.T)
dhidden1[hidden_layer1<=0.00001]=0

grads['W1']=np.dot(X.T,dhidden1)+reg*W1
grads['b1']=np.sum(dhidden1,axis=0)
```

train 函数:

加入随机成分:

```
indices=np.random.choice(num_train,batch_size,replace=True)
X_batch=X[indices]
y_batch=y[indices]
```

参数减去 学习率*梯度:

```
self.params['W1']-=learning_rate*grads['W1']
self.params['b1']-=learning_rate*grads['b1']
self.params['W2']-=learning_rate*grads['W2']
self.params['b2']-=learning_rate*grads['b2']
self.params['W3']-=learning_rate*grads['W3']
self.params['b3']-=learning_rate*grads['b3']
```

predict 函数 (取使目标函数 scores 最大的集合):

```
hidden_layer1=np.maximum(0,np.dot(X,W1)+b1)
hidden_layer2=np.maximum(0,np.dot(hidden_layer1,W2)+b2)
scores=np.dot(hidden_layer2,W3)+b3
y_pred=np.argmax(scores,axis=1)
```

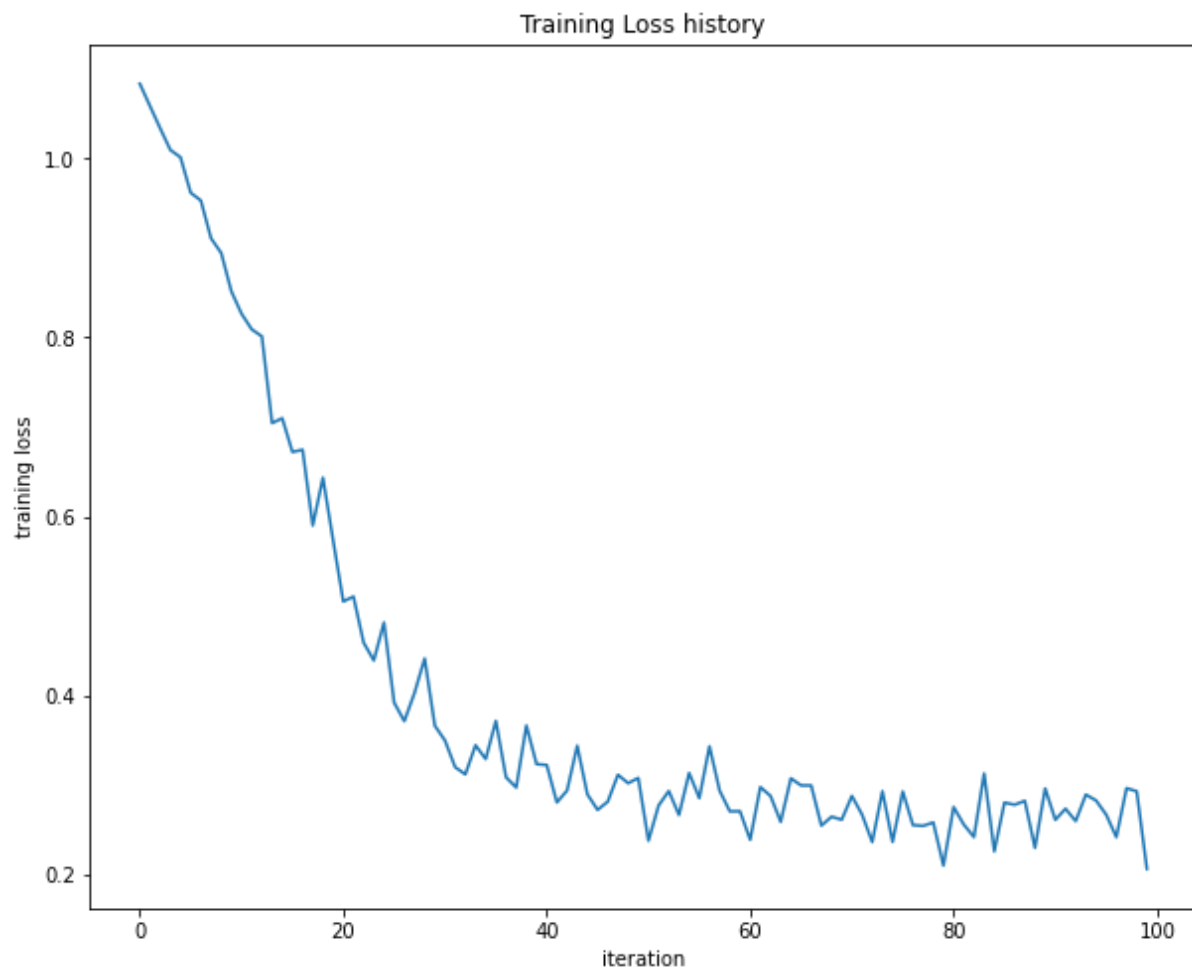

2. 完成 three_layer_net.ipynb

我的 score 和 loss 与正确值相比差异都在预期内:

```
Difference between your scores and correct scores:  
4.807962379632267e-08
```

```
Difference between your loss and correct loss:  
7.831513215705854e-13
```


训练误差迭代情况如下:



迭代 1000 次的准确率:

```
iteration 600 / 1000: loss 2.302780  
iteration 700 / 1000: loss 2.302780  
iteration 800 / 1000: loss 2.302780  
iteration 900 / 1000: loss 2.302780  
Validation accuracy: 0.113
```

loss 和 classification accuracy 的 history:



调参:

```
learning_rates=[5e-3,7e-3,9e-3]
regularization_strengths=[5e-4,7e-4,9e-3]
for lr in learning_rates:
    for rs in regularization_strengths:
        net=ThreeLayerNet(input_size,hidden_size,num_classes)
        loss_hist=net.train(X_train,y_train,X_val,y_val,
                             num_iters=10000,batch_size=150,
                             learning_rate=lr,learning_rate_decay=0.95,
                             reg=rs,verbose=False)

        val_acc = (net.predict(X_val) == y_val).mean()
        result[(lr,rs)]=val_acc
        if val_acc>best_val:
            best_val=val_acc
            best_net=net
```

模型可视化:



最终准确率及对应的 1r (0.005) 和 reg (0.0007):

```
Test accuracy: 0.528
0.005
0.0007
```

五. features

1. 完成 features.ipynb

on SVM:

调参:

```
for lr in learning_rates:
    for rs in regularization_strengths:
        svm=LinearSVM()
        loss_hist=svm.train(X_train_feats,y_train,learning_rate=lr,reg=rs,
            num_iters=1500,verbose=False)
        y_train_pred=svm.predict(X_train_feats)
        train_acc=np.mean(y_train==y_train_pred)
        y_val_pred=svm.predict(X_val_feats)
        val_acc=np.mean(y_val==y_val_pred)
        results[(lr,rs)]=(train_acc,val_acc)
        if(val_acc>best_val):
            best_val=val_acc
            best_svm=svm
```

准确率:

✓ 0.1s

0.425

on neural_net:

调参（由于运行时间太长，找到最优组合后就注释了列表）:

```
#learning_rates = [1e-2,0.1,1]
#regularization_strengths = [1e-5,1e-4,1e-3]
learning_rates = [1]
regularization_strengths = [1e-5]
for lr in learning_rates:
    for rs in regularization_strengths:
        net = ThreeLayerNet(input_dim, hidden_dim, num_classes)
        lost_hist = net.train(X_train_feats, y_train, X_val_feats, y_val,
                               num_iters=10000, batch_size=150,
                               learning_rate=lr, reg= rs, learning_rate_decay=0.95,
                               verbose=False)
        val_acc = np.mean(net.predict(X_val_feats) == y_val)
        if val_acc > best_val:
            best_val = val_acc
            best_net = net
        results[(lr,rs)] = val_acc
```

最佳组合 (lr, reg) 及准确率:

(1, 1e-05) 0.565

结论分析与体会:

1. 对比 svm 与 neural network 的参数我发现，在大多数情况下：
svm 的 lr（通常在 $1e-7$ 左右）< neural_net 的 lr（通常在 $5e-3$ 及以上）、
svm 的 reg（通常在 $1e4$ 左右）> neural_net 的 reg（通常在 $1e-4$ 及以下）、
svm 的迭代次数（通常 1500 及以下）< neural_net 的迭代次数（动辄上万），
综上，svm 似乎更好用一些（调一次 neural_net 的参数太烦人了）。

就实验过程中遇到和出现的问题，你是如何解决和处理的，自拟 1—3 道问答题:

1. 在 features.ipynb 那里，我试了很多 learning_rate 和 regularization_strength，准确率始终只有 0.1 左右（这时候迭代次数是 1500，batch_size 是 200），这让我一度怀疑自己的 neural_net.py 写错了。但是我在之前运行 three_layer_net.ipynb 的准确率达到 0.52，它同样依托 neural_net.py，证明我的 neural_net.py 是没有问题的。

```
test_acc = (best_net.predict(X_test_feats) == y_test).mean()
print(test_acc)
```

10]

Python

... 0.103

最终我将迭代次数设为 10000，batch_size 设为 150，跑出了 0.56 的准确率，最佳组合是 lr=1，reg=1e-5。

```
test_acc = (best_net.predict(X_test_feats) == y_test).mean()
print(test_acc)
```

[12] ✓ 0.1s

... 0.573

虽然取得了不错的效果，但是我总觉得舍近求远了，因为用两层的神经网络迭代 1500 次就能跑出 0.57，为什么还要用三层迭代 10000 次甚至更多呢？

```
test_acc = (best_net.predict(X_test_feats) == y_test).mean()
print(test_acc)
```

✓ 0.9s

0.561