

# 计算机科学与技术学院神经网络与深度学习课程实验报告

实验题目: Homework 2_1		学号: 201900130024
日期: 2021. 10. 11	班级: 数据 19	姓名: 刘士渤
Email: liuburger@qq.com		
<b>实验目的:</b> 掌握基本的神经网络调整技能, 并尝试改进深度神经网络:超参数调整、正则化和优化		
<b>实验软件和硬件环境:</b> VScode JupyterNoteBook 联想拯救者 Y7000p		
<b>实验原理和方法:</b> neural network		
<b>实验步骤: (不要求罗列完整源代码)</b> 1. 补全 Initialization.ipynb 根据 $W_l$ 的形状是 $(layers\_dims[L], layers\_dims[L-1])$ $b_l$ 的形状是 $(layers\_dims[L], 1)$ 得代码: <pre>parameters['W' + str(l)] = np.zeros([layers_dims[l], layers_dims[l-1]]) parameters['b' + str(l)] = np.zeros([layers_dims[l], 1])</pre> 根据随机的要求得代码: <pre>parameters['W' + str(l)] = np.random.randn(layers_dims[l], layers_dims[l-1])*10 parameters['b' + str(l)] = np.zeros([layers_dims[l], 1])</pre> 根据提示开方得代码: <pre>parameters['W' + str(l)] = np.random.randn(layers_dims[l], layers_dims[l-1])*np.sqrt(2./layers_dims[l-1]) parameters['b' + str(l)] = np.zeros([layers_dims[l], 1])</pre> 2. 补全 Gradient Checking.ipynb 根据 $J(\theta) = \theta * x$ 得代码: <pre>J = theta*x</pre> 根据 $d\theta = \frac{\partial J}{\partial \theta} = x.$ 得代码: <pre>dtheta = x</pre>		

根据

1.  $\theta^+ = \theta + \varepsilon$
2.  $\theta^- = \theta - \varepsilon$
3.  $J^+ = J(\theta^+)$
4.  $J^- = J(\theta^-)$
5.  $gradapprox = \frac{J^+ - J^-}{2\varepsilon}$

$$difference = \frac{\|grad - gradapprox\|_2}{\|grad\|_2 + \|gradapprox\|_2}$$

得代码:

```
thetaplus = theta+epsilon # Step
thetaminus = theta-epsilon # Step
J_plus = thetaplus*x # Step
J_minus = thetaminus*x # Step
gradapprox = (J_plus-J_minus)/(2*epsilon) # Step
### END CODE HERE ###

# Check if gradapprox is close enough to the output of backward
### START CODE HERE ### (approx. 1 line)
grad = x
### END CODE HERE ###

### START CODE HERE ### (approx. 1 line)
numerator = np.linalg.norm(grad-gradapprox)
denominator = np.linalg.norm(grad)+np.linalg.norm(gradapprox)
difference = numerator/denominator
```

根据

- To compute `J_plus[i]`:
  1. Set  $\theta^+$  to `np.copy(parameters_values)`
  2. Set  $\theta_i^+$  to  $\theta_i^+ + \varepsilon$
  3. Calculate  $J_i^+$  using `forward_propagation_n(x, y, vector_to_dictionary( $\theta^+$ ))`.
- To compute `J_minus[i]`: do the same thing with  $\theta^-$
- Compute  $gradapprox[i] = \frac{J_i^+ - J_i^-}{2\varepsilon}$

$$difference = \frac{\|grad - gradapprox\|_2}{\|grad\|_2 + \|gradapprox\|_2}$$

得代码:

```
thetaplus = np.copy(parameters_values)
thetaplus[i][0] +=epsilon # Step 2
J_plus[i], _ = forward_propagation_n(X,Y,vector_to_dictionary(thetaplus))
### END CODE HERE ###

# Compute J_minus[i]. Inputs: "parameters_values, epsilon". Output = "J_minus[i]"
### START CODE HERE ### (approx. 3 lines)
thetaminus = np.copy(parameters_values)
thetaminus[i][0] -=epsilon # Step 2
J_minus[i], _ = forward_propagation_n(X,Y,vector_to_dictionary(thetaminus))
### END CODE HERE ###

# Compute gradapprox[i]
### START CODE HERE ### (approx. 1 line)
gradapprox[i] = (J_plus[i]-J_minus[i])/(2*epsilon)
### END CODE HERE ###

# Compare gradapprox to backward propagation gradients by computing difference.
### START CODE HERE ### (approx. 1 line)
numerator = np.linalg.norm(grad-gradapprox) # Step 1
denominator = np.linalg.norm(grad)+np.linalg.norm(gradapprox) # Step 2
difference = numerator/denominator # Step 3
```

### 3. 补全 Optimization methods.ipynb

根据

$$W^{[l]} = W^{[l]} - \alpha dW^{[l]}$$
$$b^{[l]} = b^{[l]} - \alpha db^{[l]}$$

得代码:

```
parameters["W" + str(l+1)] -= learning_rate*grads['dW' + str(l+1)]
parameters["b" + str(l+1)] -= learning_rate*grads['db' + str(l+1)]
```

根据

```
first_mini_batch_X = shuffled_X[:, 0 : mini_batch_size]
second_mini_batch_X = shuffled_X[:, mini_batch_size : 2 * mini_batch_size]
```

得代码:

```
mini_batch_X = shuffled_X[:,k*mini_batch_size:(k+1)*mini_batch_size]
mini_batch_Y = shuffled_Y[:,k*mini_batch_size:(k+1)*mini_batch_size]
```

根据

$$(m - \text{mini\_batch\_size} \times \lfloor \frac{m}{\text{mini\_batch\_size}} \rfloor).$$

得代码:

```
mini_batch_X = shuffled_X[:,num_complete_minibatches*mini_batch_size:m]
mini_batch_Y = shuffled_Y[:,num_complete_minibatches*mini_batch_size:m]
```

根据

```
v["dW" + str(l+1)] = ... #(numpy array of zeros with the same shape as parameters["W" + str(l+1)])
v["db" + str(l+1)] = ... #(numpy array of zeros with the same shape as parameters["b" + str(l+1)])
```

得代码:

```
v["dW" + str(l+1)] = np.zeros([parameters['W' + str(l+1)].shape[0], parameters['W' + str(l+1)].shape[1]])
v["db" + str(l+1)] = np.zeros([parameters['b' + str(l+1)].shape[0], parameters['b' + str(l+1)].shape[1]])
```

根据

$$\begin{cases} v_{dW^{[l]}} = \beta v_{dW^{[l]}} + (1 - \beta) dW^{[l]} \\ W^{[l]} = W^{[l]} - \alpha v_{dW^{[l]}} \\ \\ \begin{cases} v_{db^{[l]}} = \beta v_{db^{[l]}} + (1 - \beta) db^{[l]} \\ b^{[l]} = b^{[l]} - \alpha v_{db^{[l]}} \end{cases} \end{cases}$$

得代码:

```
# compute velocities
v["dW" + str(l+1)] = beta*v["dW" + str(l+1)]+(1-beta)*grads['dW' + str(l+1)]
v["db" + str(l+1)] = beta*v["db" + str(l+1)]+(1-beta)*grads['db' + str(l+1)]
# update parameters
parameters["W" + str(l+1)] -= learning_rate*v["dW" + str(l+1)]
parameters["b" + str(l+1)] -= learning_rate*v["db" + str(l+1)]
```

根据

```
v["dW" + str(l+1)] = ... #(numpy array of zeros with the same shape as parameters["W" + str(l+1)])
v["db" + str(l+1)] = ... #(numpy array of zeros with the same shape as parameters["b" + str(l+1)])
s["dW" + str(l+1)] = ... #(numpy array of zeros with the same shape as parameters["W" + str(l+1)])
s["db" + str(l+1)] = ... #(numpy array of zeros with the same shape as parameters["b" + str(l+1)])
```

得代码:

```
v["dW" + str(l+1)] = np.zeros([parameters['W' + str(l+1)].shape[0], parameters['W' + str(l+1)].shape[1]])
v["db" + str(l+1)] = np.zeros([parameters['b' + str(l+1)].shape[0], parameters['b' + str(l+1)].shape[1]])
s["dW" + str(l+1)] = np.zeros([parameters['W' + str(l+1)].shape[0], parameters['W' + str(l+1)].shape[1]])
s["db" + str(l+1)] = np.zeros([parameters['b' + str(l+1)].shape[0], parameters['b' + str(l+1)].shape[1]])
```

根据

$$\begin{cases} v_{W^{[l]}} = \beta_1 v_{W^{[l]}} + (1 - \beta_1) \frac{\partial J}{\partial W^{[l]}} \\ v_{W^{[l]}}^{corrected} = \frac{v_{W^{[l]}}}{1 - (\beta_1)^t} \\ s_{W^{[l]}} = \beta_2 s_{W^{[l]}} + (1 - \beta_2) \left( \frac{\partial J}{\partial W^{[l]}} \right)^2 \\ s_{W^{[l]}}^{corrected} = \frac{s_{W^{[l]}}}{1 - (\beta_2)^t} \\ W^{[l]} = W^{[l]} - \alpha \frac{v_{W^{[l]}}^{corrected}}{\sqrt{s_{W^{[l]}}^{corrected} + \epsilon}} \end{cases}$$

## 得代码：

```
v["dW" + str(l+1)] = beta1*v["dW" + str(l+1)]+(1-beta1)*grads['dW' + str(l+1)]
v["db" + str(l+1)] = beta1*v["db" + str(l+1)]+(1-beta1)*grads['db' + str(l+1)]
#### END CODE HERE ####

# Compute bias-corrected first moment estimate. Inputs: "v, beta1, t". Output: "v_corrected".
#### START CODE HERE #### (approx. 2 lines)
v_corrected["dW" + str(l+1)] = v["dW" + str(l+1)]/(1-beta1**t)
v_corrected["db" + str(l+1)] = v["db" + str(l+1)]/(1-beta1**t)
#### END CODE HERE ####

# Moving average of the squared gradients. Inputs: "s, grads, beta2". Output: "s".
#### START CODE HERE #### (approx. 2 lines)
s["dW" + str(l+1)] = beta2*s["dW" + str(l+1)]+(1-beta2)*(grads['dW' + str(l+1)]**2)
s["db" + str(l+1)] = beta2*s["db" + str(l+1)]+(1-beta2)*(grads['db' + str(l+1)]**2)
#### END CODE HERE ####


# Compute bias-corrected second raw moment estimate. Inputs: "s, beta2, t". Output: "s_corrected".
#### START CODE HERE #### (approx. 2 lines)
s_corrected["dW" + str(l+1)] = s["dW" + str(l+1)]/(1-beta2**t)
s_corrected["db" + str(l+1)] = s["db" + str(l+1)]/(1-beta2**t)
#### END CODE HERE ####

# Update parameters. Inputs: "parameters, learning_rate, v_corrected, s_corrected, epsilon". Output: "parameters".
#### START CODE HERE #### (approx. 2 lines)
parameters["w" + str(l+1)] -= learning_rate*v_corrected["dW" + str(l+1)]/(np.sqrt(s_corrected["dW" + str(l+1)]))+epsilon
parameters["b" + str(l+1)] -= learning_rate*v_corrected["db" + str(l+1)]/(np.sqrt(s_corrected["db" + str(l+1)]))+epsilon
```


## 结论分析与体会：

### 1. Initialization


原数据：




采用全 0 初始化的 cost 与 iterations 的关系（因为初始化为 0 所以迭代没有任何效果）：




预测准确率为 0.5:




采用随机数初始化的 cost 与 iterations 的关系（损失函数随迭代次数增加而降低，速度为快→慢→快→慢）：




预测准确率为 0.83



采用 he 初始化的 cost 与 iterations 的关系（损失函数随迭代次数增加而降低，速度同样为快→慢→快→慢，但较随机初始化更为平滑）：



预测准确率为 0.99





## 2. Gradient Checking

写了梯度估计后，与梯度做差，发现差异很小：

```
The gradient is correct!  
difference = 2.919335883291695e-10
```

修改了错误代码之后


```
dw2 = 1./m * np.dot(dZ2, A1.T)  
db2 = 1./m * np.sum(dZ2, axis=1, keepdims = True)  
  
dA1 = np.dot(W2.T, dZ2)  
dZ1 = np.multiply(dA1, np.int64(A1 > 0))  
dw1 = 1./m * np.dot(dZ1, X.T)  
db1 = 1./m * np.sum(dZ1, axis=1, keepdims = True)
```

差异由 0.29 变为  $1.19e-7$


```
Your backward propagation works perfectly fine! difference = 1.1890913024229996e-07
```

## 3. Optimization methods


原数据




采用梯度下降方法的 cost 与 epochs 的关系（损失函数随迭代次数增加而降低，总体呈下降趋势但还没有趋于定值）：




预测准确率为 0.79




采用动量方法的 cost 与 epochs 的关系（图像似乎和梯度下降的一模一样）：



预测准确率为 0.79



采用动量方法的 cost 与 epochs 的关系（损失函数随迭代次数增加而降低，最终稳定在 0.15 左右）：



预测准确率为 0.94





就实验过程中遇到和出现的问题，你是如何解决和处理的，自拟 1—3 道问答题：



1. Initialization 部分，我尝试着提升 random 初始化的准确率。在将 W 的系数从 10 改为 1.7 之后

```
parameters['W' + str(l)] = np.random.randn(layers_dims[l], layers_dims[l-1])*1.7
```

损失函数随迭代较之前收敛更快：





预测准确率也提高到了 0.99





只是改了一下系数，却达到了 he 方法的准确率；看来 1.7 这个数值和对 2/前一层的维度 开方是差不多的。



2. Optimization 部分，由于观察到采用梯度下降和动量方法的损失函数曲线几乎一模一样，并且波动较大，趋于定值的趋势不明显，所以我决定提高 learning rate。




将 learning rate 提高 100 倍即 0.07 后，损失函数就像 Adam 一样收敛了。



预测准确率提高到了 0.94



但是 Adam 的损失函数却波动起来了



且准确率也降低到 0.87，看来梯度下降和动量方法适合较大的学习率，而 Adam 方法适合较小的学习率。

